

In Class Discussion

EN 601.124: The Ethics of Artificial Intelligence and Automation

Department of Computer Science, February 11th 2026



Moral Machine





What Moral machine findings were you most surprised about? List 3-5.



The Moral Machine found ethics are not uniform across cultures.

How can we integrate this fact into AI systems so they uphold the appropriate situational ethics?

If values differ this much across societies, how could an ethical framework ground the design of globally deployed AI?




In the lecture we discussed how the utilitarian framework has a "measurement problem" as welfare and harm are hard to quantify.

The Moral Machine circumvents it—it presents ethical dilemmas as forced binary choices (e.g., swerve left or right). How does this framing compare to the hospital AI scenario from class? What gets lost when we reduce ethical reasoning to discrete choices rather than open-ended deliberation?

Alexandria





What makes something a value system?
Are these things random/justified?
Look up and read about Schwartz Theory of Basic Values.
Discuss if this framework can serve as a foundation instead.



Alexandria frames value pluralism as a design goal — letting different users prioritize different things.

How would a critical theorist evaluate this approach?

Does offering individual value customization obscure structural power dynamics, like who built the value library or which values were included and excluded?

Or does it instead help rebalance the power dynamics?





**An abolitionist framework asks
"should this system exist at all?"**

**Is it reasonable to question whether
algorithmically re-ranked feeds should
exist (why not just chronological
feed), rather than optimizing which
values they serve?**

Synthesis





What decisions / control of the AI model should users be able to make and which should they be unable to make? Why?



What decisions / control of the AI model should users be able to make and which should they be unable to make? Why?

- **Alexandria implements values as LLM-powered classifiers that re-rank content in real time. From a deontological perspective, are there certain values that should never be available as options — regardless of user preference — because offering them would violate a moral duty?**
- **The Moral Machine found that people in some clusters preferred saving higher-status individuals. If we take these preferences seriously as data, how should engineers decide which elicited values to encode and which to override?**



Think of three AI systems that you use in your everyday life. What would better transparency of philosophical foundation for this system entail?



What can alexandria teach us about ethical design of AI chatbots in general, beyond social media?

How can pluralism be generalized to other AI systems in industries where failure has more severe consequences than social media?