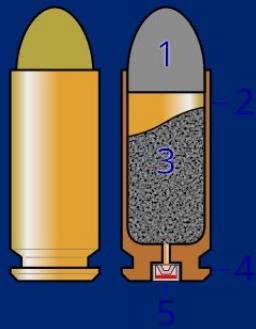


# Can KV-cache serve as compact memory module when reasoning over long or unbounded contexts?

Ziyang Huang, Yingfei Xu



# Cartridges: Lightweight and general-purpose long context representations via self-study

**Sabri Eyuboglu, Ryan Saul Ehrlich, Simran Arora, Neel Guha,  
Dylan Zinsley, Emily Ruoyu Liu, Atri Rudra, James Y. Zou,  
Azalia Mirhoseini, Christopher Re**



# Long-context ICL is accurate but expensive: $KV \propto$ input length.

---

But can we use a swappable KV cache module and reuse it?

## Cartridges

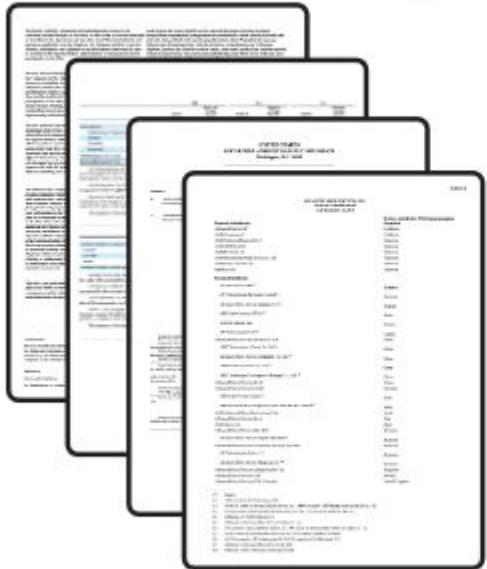
Turning long corpora into tiny, reusable KV prefixes for fast, high-quality answers.

## Performance

ICL-like quality with dramatically less memory and higher throughput when many queries target the same corpus.

## Problem Setting

Document Corpus



Queries

Please, summarize  
the documents...



Write a rock song  
about the docs...



What is the D&A  
margin for FY15...



*Users send **many** messages grounded  
in a **single** large corpus of text.*

# Motivation

## In-context learning



Prefill

KV Cache

k[1] v[1]

Documents represented by KV cache produced with standard **prefill**.

- ✓ Supports general queries
- ✗ High GPU memory consumption

k[2] v[2]

k[3] v[3]

k[4] v[4]

LLM + KV Cache

What is the D&A margin for FY15...

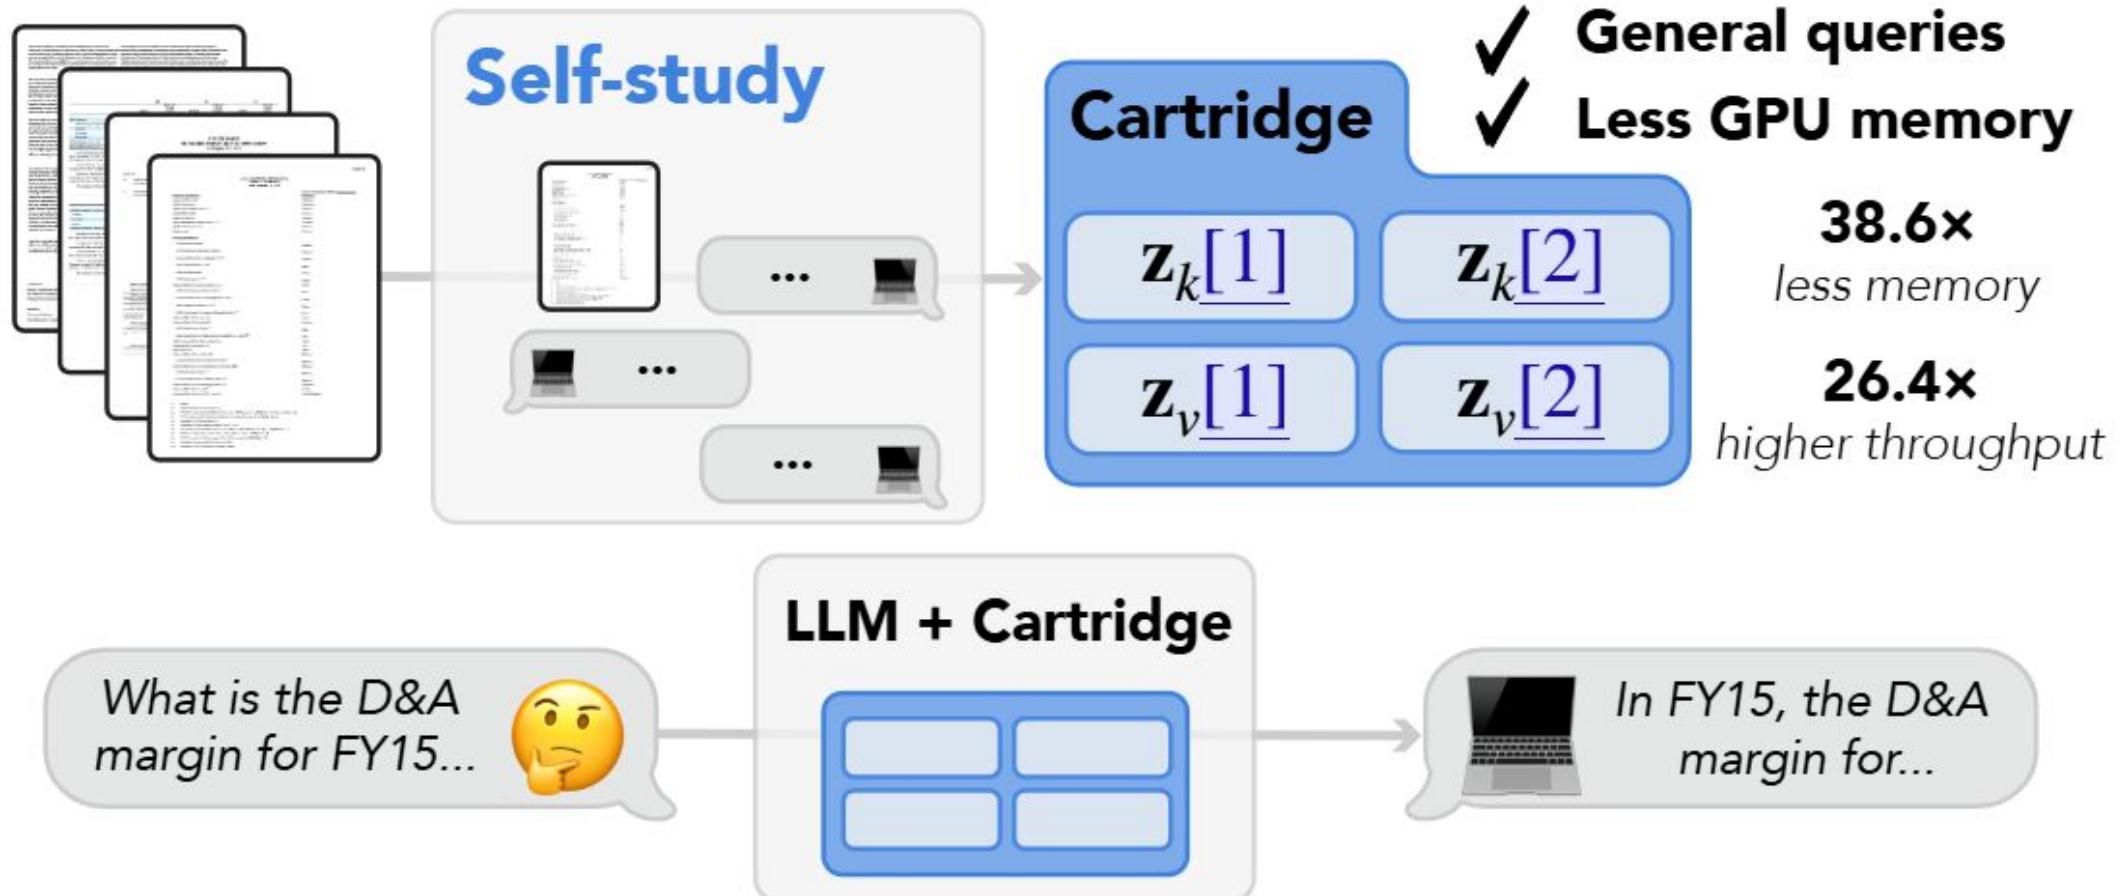


In FY15, the D&A margin for...

# Motivation

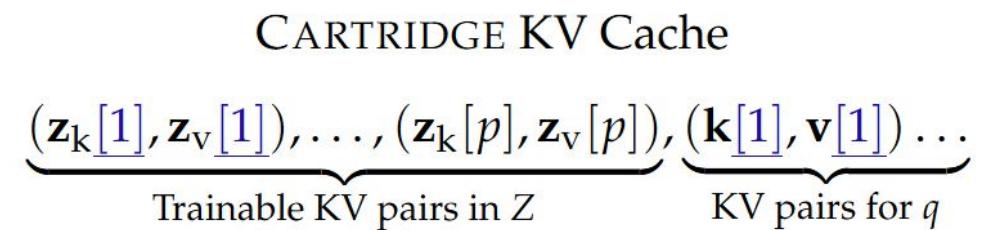
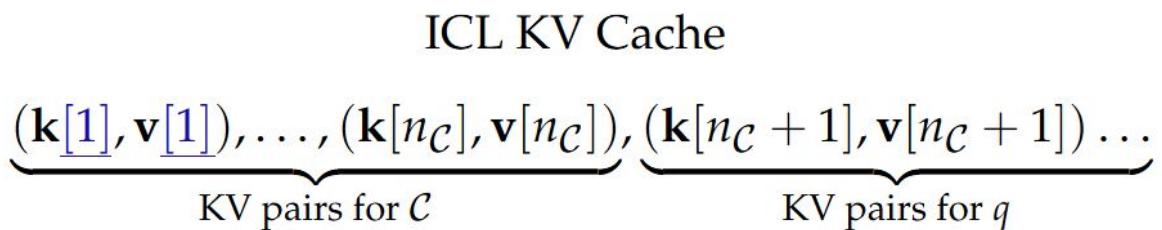
## Cartridges

Documents represented with a compressed KV cache  
that is trained with ***self-study***.



# What is a Cartridge

- A learned KV prefix of length  $p$  (e.g., 128–8192 tokens) attached at inference time.
- No base-model weight updates; just load/unload per corpus.
- A composable module; concatenate multiple cartridges (e.g., Policy  $\oplus$  API Docs  $\oplus$  Codebase) for cross-source reasoning.



Analogy: Like prefix-tuning, but trained to imitate ICL on that corpus.



# Self-study Training

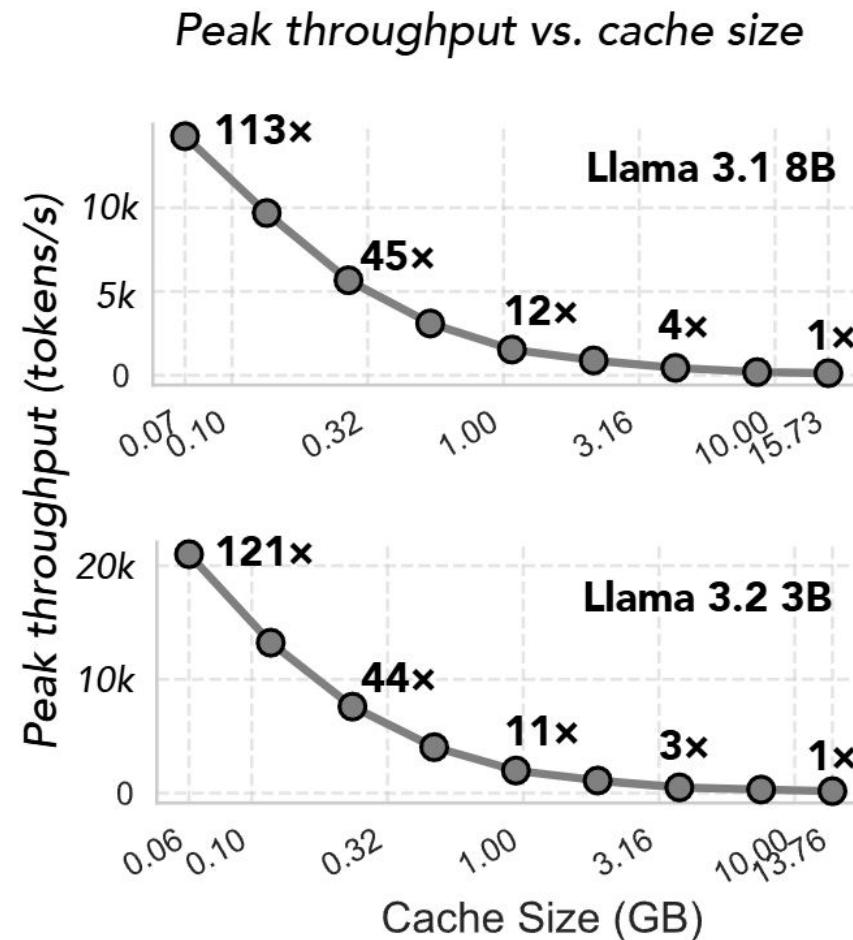
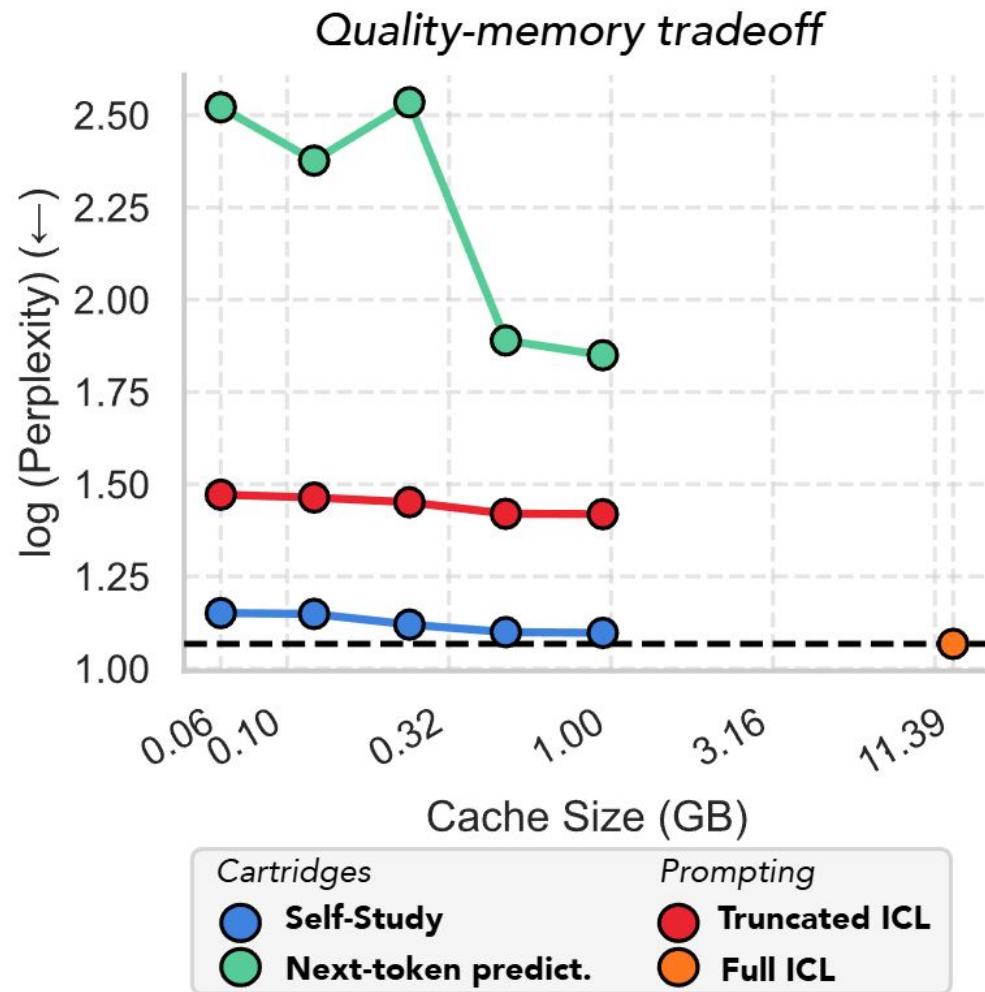
---

- Synthesize dialogs about the corpus: For each chunk  $\tilde{c}$ , auto-generate Q $\leftrightarrow$ A / instruction traces.
- Teacher = base LLM with  $\tilde{c}$  in context; Student = same LLM + trainable cartridge Z without  $\tilde{c}$ .
- Objective: minimize step-wise  $\text{KL}(p_{\text{t}} \parallel p_{\text{s}})$  over next-token distributions.

Intuition: Student learns to act as if the corpus were present.

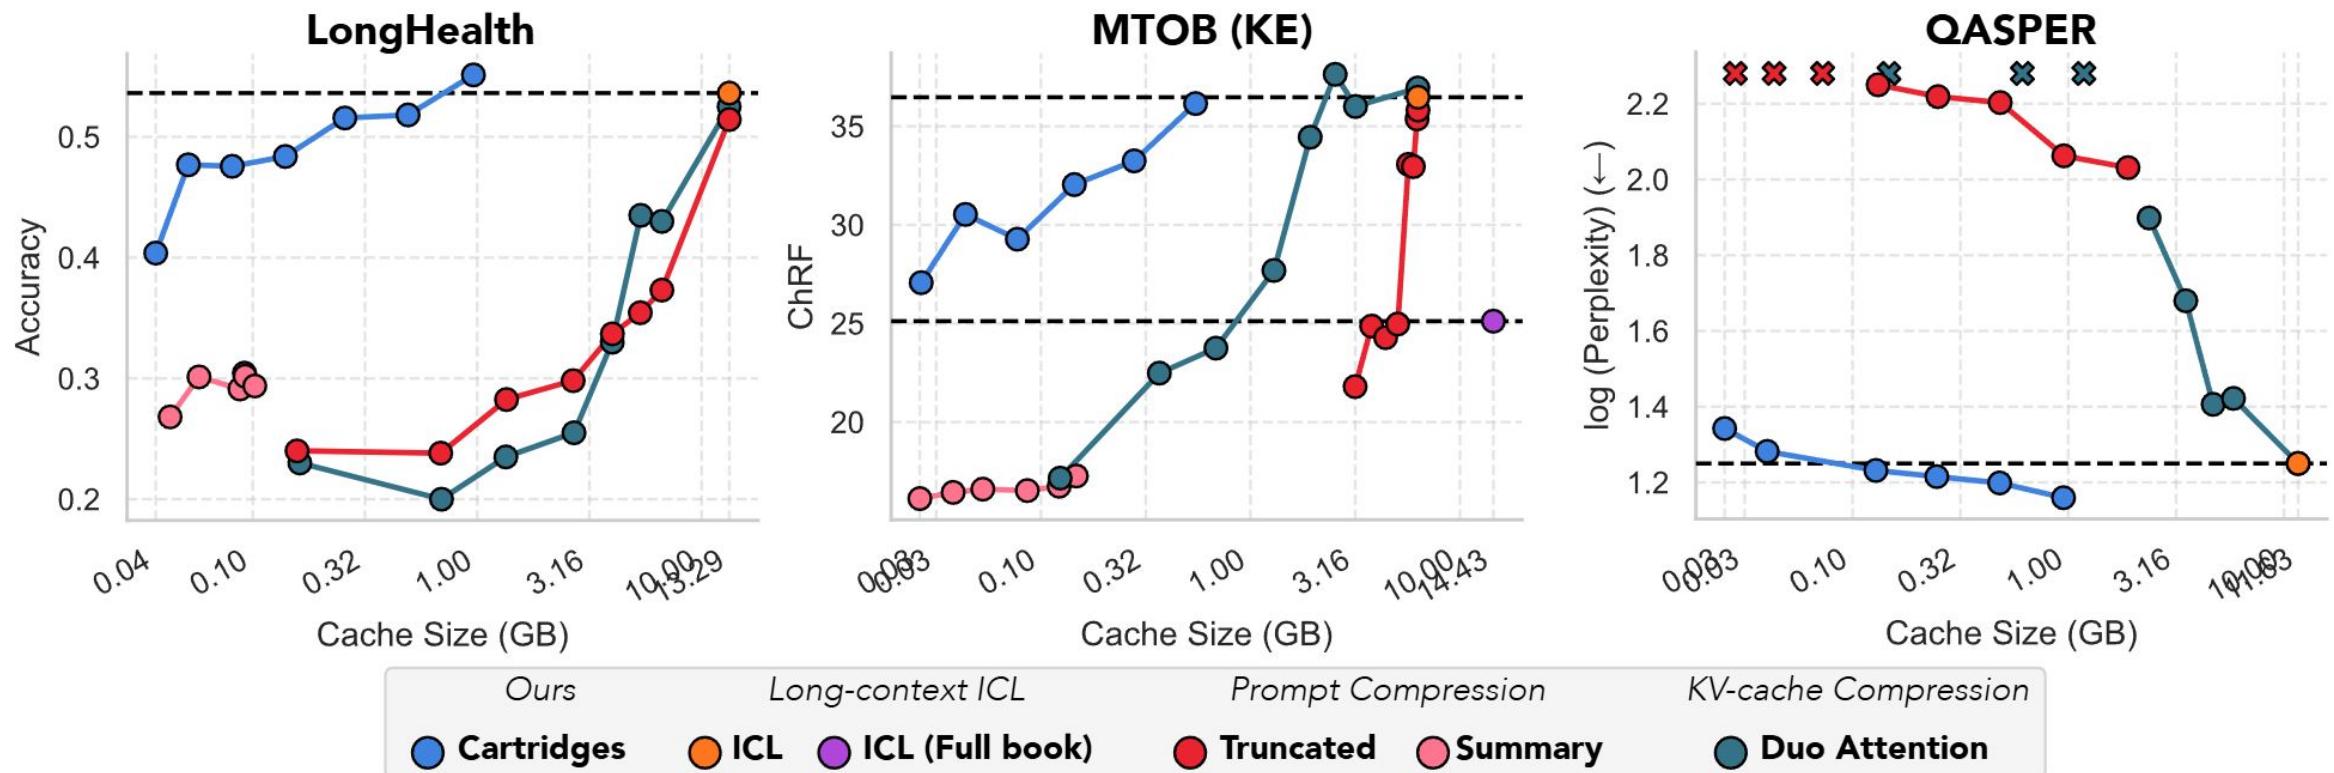


# Performance: throughput and cache size



# Performance: Cartridges vs baselines

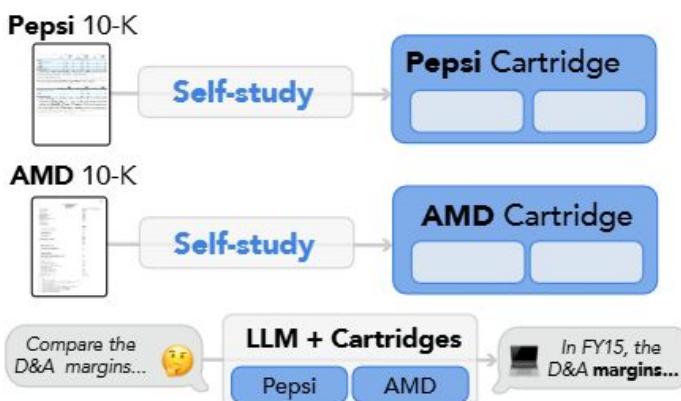
Cartridges matches ICL quality with lower memory costs.



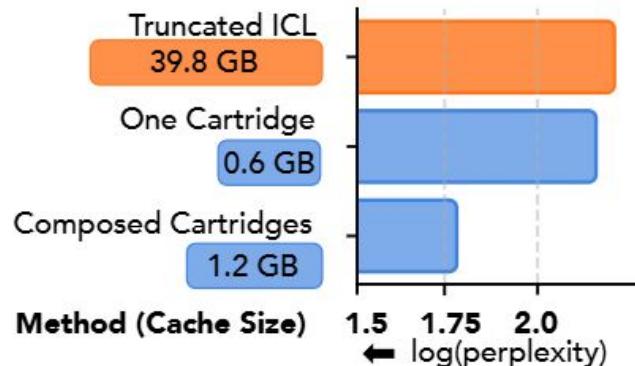
# Surprising notes

- Context stretch: Effective context extended (e.g., 128k → ~484k on MTOB)
- Composition: Multiple cartridges combine without re-training.

## Cartridge Composition



## Multi-doc Question Answering



Who audited the Boeing and AMD statements, respectively? 🤔

The audit of the consolidated financial statements of AMD was performed by Ernst & Young LLP, while the audit of the consolidated financial statements of Boeing was performed by Deloitte & Touche LLP.

List a few competitors for each of PepsiCo and AMD as stated in each 10K. 😊

Here are some competitors for PepsiCo and AMD:  
\* Unilever (as a competitor) ...  
\* Red Bull (as a competitor in the energy drink market)  
AMD:  
\* Intel (as a competitor in the ...



# Cartridges vs RAG? LoRA?

---

RAG Good for flexible and live updates, but retrieval and long prompts are still costly.

-> Cartridges are optimized on stable, repeatedly-queried corpora.

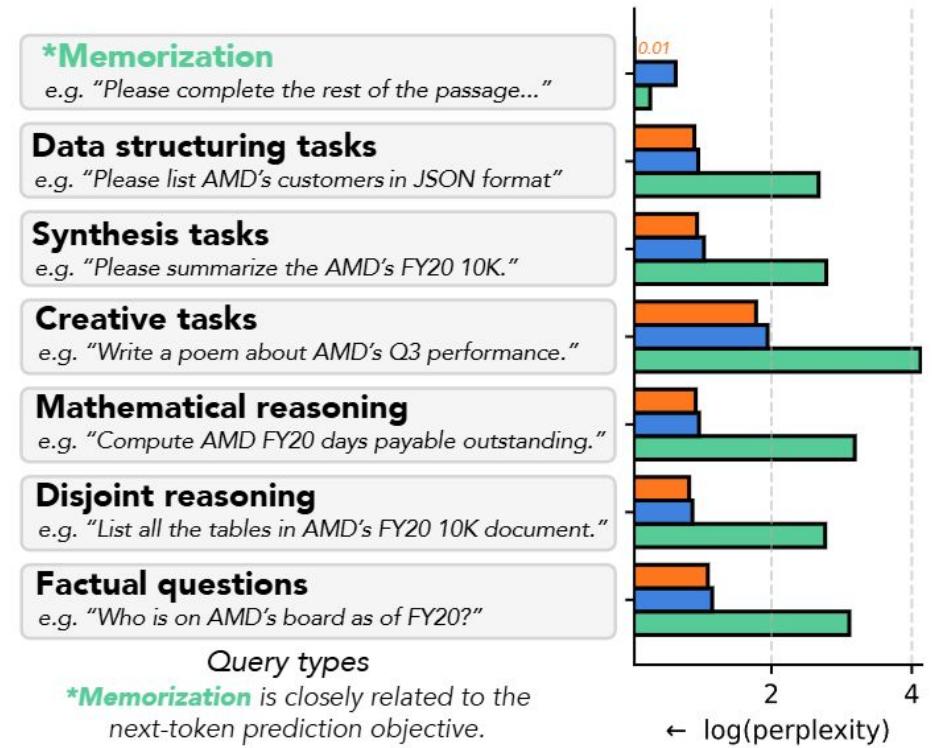
LoRA Much more expressive for training, but also to serve on infrastructures.

-> Cartridges are prefix-only; also outperforms LoRA.

# Other discussions

- A Cartridge can be served efficiently with minimal changes to existing LLM inference servers (e.g. SGLang)
- Limitations:
  - Upfront compute: Strong performance but training must be amortized
  - Coverage: training rely on synthetic dialogs; bad dialogs cause blind spots
  - Timeliness: Corpus changes require re-training or incremental fine-tuning

Generalization to diverse queries



## Cartridges

● Self-Study

● Next-token predict.

## Prompting

● Truncated ICL

● Full ICL



# Takeaways

---

Cartridges = compact, reusable memory of a corpus with ICL-like behavior.

Use Self-study (synthetic dialogs + context distillation) for training.

For repeated queries, strong performance in terms of larger throughput and reduced memory

# **BumbleBee: Dynamic KV-Cache Streaming Submodular Summarization for Infinite-Context Transformers**

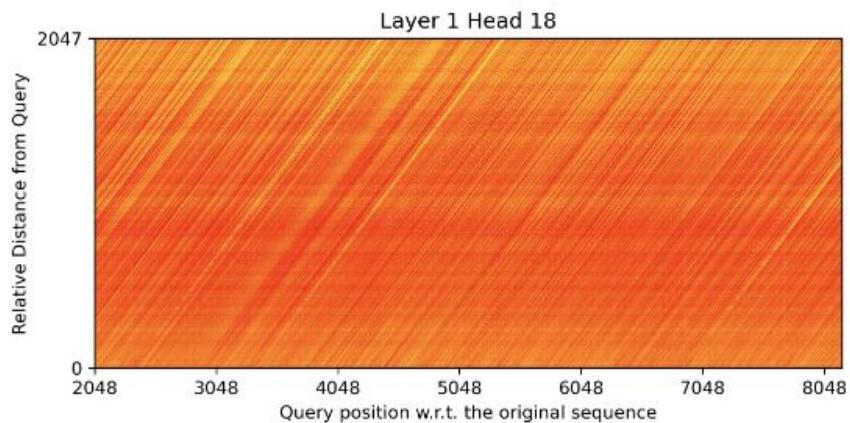
**Lilly Kumari, Shengjie Wang, Tianyi Zhou, Nikhil Sarda,  
Anthony Rowe, Jeff Bilmes**



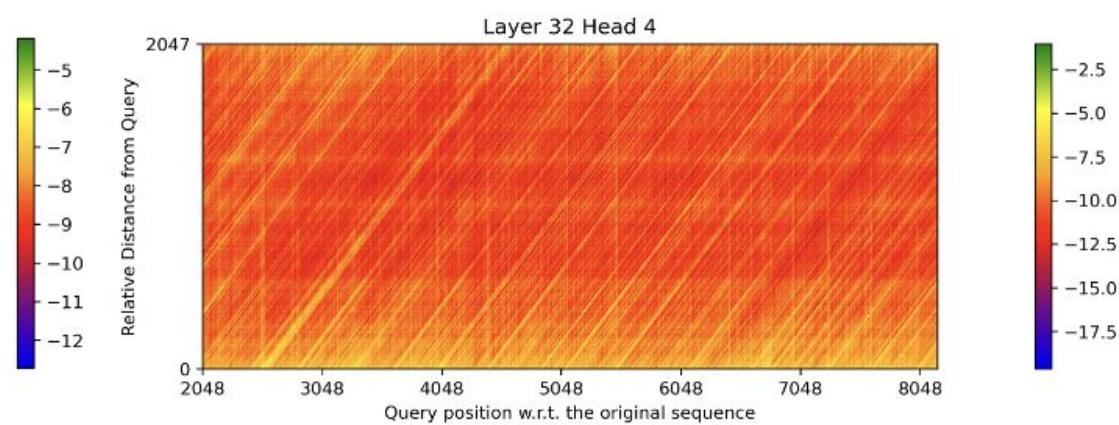
# Motivation

---

strongly attends to a small subset of tokens



(a) Test sample 1



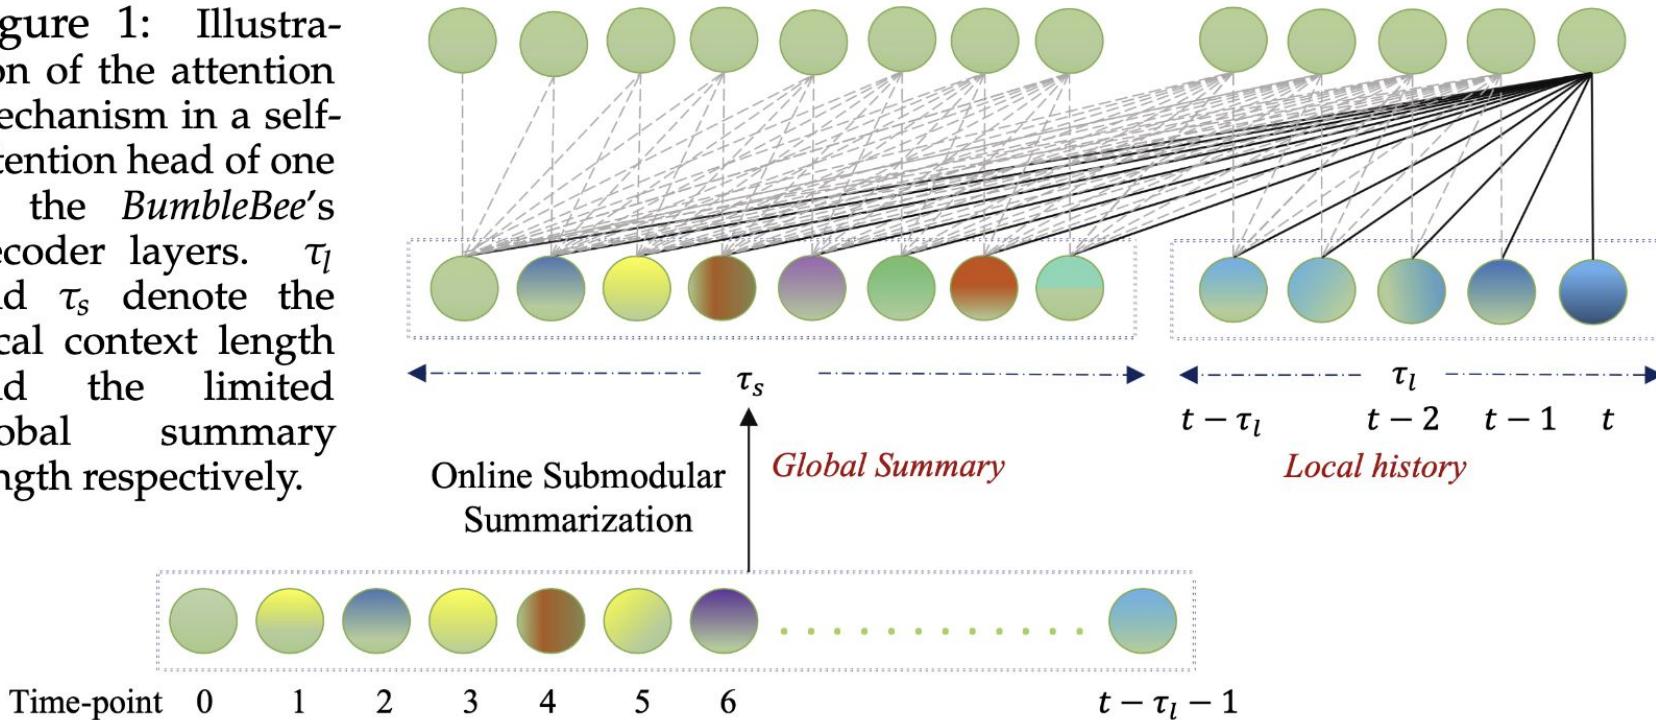
(b) Test sample 2

Figure 2: Attention maps for two different WikiText-103 articles using LLaMA-7B model.

# Core Idea

- Existing methods use **modular scores**, where each KV state is evaluated independently
- Frame the KV cache selection as **subset selection** problem with submodular objective function

Figure 1: Illustration of the attention mechanism in a self-attention head of one of the *BumbleBee*'s decoder layers.  $\tau_l$  and  $\tau_s$  denote the local context length and the limited global summary length respectively.



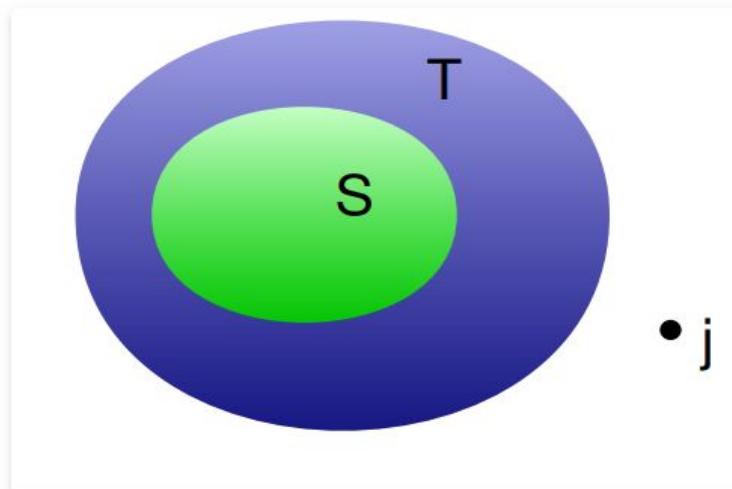
# Submodular

---

A submodular function  $f : 2^V \rightarrow R$  defined on ground set  $V$ , it has **diminishing return property**:

$$\text{if } S \subset T, j \notin T, f(S \cup \{j\}) - f(S) \geq f(T \cup \{j\}) - f(T)$$

Greedy algorithm for submodular maximization has theoretical  $(1 - 1/e)$  guarantee from the optimal solution



# Submodularity

## facility location (FL) function

how well a subset can represent the whole set

$$f_{\text{FL}}(A) = \sum_{v \in V} \max_{v' \in A} \text{sim}(v, v'). \quad (2)$$

## feature-based function

how much total importance is captured

$$c(A) = \sum_{u \in U} \phi_u(\sum_{v \in A} m_u(v)). \quad (3)$$

How to select a subset?

Diversity

Importance

# Submodularity

---

## For Bumblebee

- $c(\cdot)$ :  $|U| = 1$  and weight  $m_u(k_i) = a_n^i$ , accumulated attention
- $f_{FL}(\cdot)$ : similarity matrix with pairwise cosine, followed by ReLU

## Mixture function: trade-off

normalized with  $f_{FL}(\emptyset) = 0$  and  $f_{FL}(V) = 1$  (and the same for  $c(\cdot)$ )

$$g_\lambda(A) = \lambda f_{FL}(A) + (1 - \lambda)c(A)$$



# Offline algorithm

---

---

**Algorithm 1** Offline Submodular KV cache Summarization during Prefill/Encoding Phase

---

- 1: **Input:** Submodular functions capturing diversity  $f_{\text{FL}}$  in the key embeddings space and importance  $c$  via attention frequency for layer  $l$  and attention head  $h$ ; mixture function  $g_\lambda(\cdot) = \lambda f_{\text{FL}}(\cdot) + (1 - \lambda)c(\cdot)$ ; a set of  $n$  KV attention states  $K_n = \{(k_i)\}_{i=1}^n, V_n = \{(v_i)\}_{i=1}^n$  corresponding to the  $n$  prompt tokens; budget  $\tau_s$ .
  - 2: **Output:** A final summary  $S_n$  such that  $S_n \subseteq \{(k_i, v_i)\}_{i=1}^n$  and  $|S_n| \leq \tau_s$ .
  - 3: **Initialize:**  $S_n = \emptyset$ ; compute accumulated attention score vectors  $a_n$  for each key  $k \in \{k_i\}_{i=1}^n$ .  $a_n^i$  denotes accumulated attention scores attributed to key  $k_i$  across all  $n$  query tokens.
  - 4: **for**  $j = 1$  to  $\tau_s$  **do**
  - 5:    $k_{\text{imp}} \leftarrow \operatorname{argmax}_{e \in K_n \setminus S_n} g_\lambda(S_n \cup e) - g(S_n)$
  - 6:    $S_n \leftarrow S_n \cup \{(k_{\text{imp}}, v_{\text{imp}})\}$  where  $v_{\text{imp}}$  is the value embedding associated with  $k_{\text{imp}}$ .
  - 7: **end for**
-

# Online algorithm

fill until budget,  
else greedy choose  
worst one to evict

for each new token  
compute costs

$$\mathcal{O}(\tau_s \times d + \tau_s^2)$$

---

**Algorithm 2** *BumbleBee: Streaming Submodular KV cache Summarization for Transformers*

---

- 1: **Input:** Submodular functions for diversity  $f_{FL}$  in the key embeddings space and importance  $c$  w.r.t. attention frequency resp. for layer  $l$  and attention head  $h$ ; mixture function  $g_\lambda(\cdot) = \lambda f_{FL}(\cdot) + (1 - \lambda)c(\cdot)$ ; stream of QKV attention states  $\{(q_i, k_i, v_i)\}_{i=1}^n$ ; budget  $\tau_s$ .
- 2: **Output:** A running summary  $S_t$  of for every time step  $t$  such that  $S_t \subseteq \{(k_i, v_i)\}_{i=1}^t$ .
- 3: **Initialize:**  $S_0 = \emptyset, a_0 = \emptyset$  where  $a_t \in \mathbf{R}^{|S_t|}$  denotes the accumulated attention scores corresponding to keys present in  $S_t$  across  $t$  time steps.
- 4: **for**  $t = 1, \dots, n$  **do**
- 5:   Update  $a_t$  for each  $k \in S_{t-1}$  by adding  $a(q_t, k, S_{t-1} \cup k_t)$
- 6:   **if**  $t < \tau_s$  **then**
- 7:      $S_t \leftarrow S_{t-1} \cup \{(k_t, v_t)\}$
- 8:     Append  $a(q_t, k_t, S_t)$  to  $a_t$  s.t.  $|a_t| = |S_t|$
- 9:   **else**
- 10:     Let  $S'_t = S_{t-1} \cup \{(k_t, v_t)\}; k_{\text{discard}} \leftarrow \operatorname{argmin}_{k_i \in S'_t} g_\lambda(k_i | S'_t \setminus k_i)$
- 11:      $S_t \leftarrow S'_t \setminus \{(k_{\text{discard}}, v_{\text{discard}})\}$
- 12:     **if**  $k_{\text{discard}} \neq k_t$  **then**
- 13:       Evict  $a_t^j$  (the accumulated attention score for the discarded key  $k_{\text{discard}}$ ) from  $a_t$ .
- 14:       Append  $a(q_t, k_t, S_t)$  to  $a_t$
- 15:     **end if**
- 16:   **end if**
- 17: **end for**

# Experiments

---

- **Datasets from benchmark:** Im-eval-harness, HELM, LongBench
- **Models:** LLaMA 7B and 13B, LLaMA2 7B and 13B, Llama-2-Chat 7B and LongChat-32k 7B
- **Baselines:** All, Local, Random+Local, Attention sinks+Local, H2+local
- ***Submarine*** software system for submodular computation

# Experiments

Ilm-eval-harness benchmark, **0.1x** the input length as budget

( log-based  $\phi(x) = \log(1 + x)$ )

( power-based  $\phi(x) = g^{-1}(x)$  where  $g(y) = \alpha y^{1/\alpha} + \beta y$ )

Model	Methods	OpenBookQA	COPA	RTE	MathQA	PiQA	Winogrande
LLaMA-13B	All	47.4	85	73.28	31.86	80.36	75.69
	Local	28.4	64	53.43	23.25	58.32	49.88
	Random + Local	27.6	58	54.63	21.76	54.13	50.64
	Attn Sinks + Local	44.4	80	67.51	29.78	79.22	70.48
	H2 + Local	44.2	83	64.98	29.71	<b>79.49</b>	70.32
	BumbleBee	<b>47.6</b>	<b>85</b>	<b>71.48</b>	<b>31.02</b>	79.38	71.98
	BumbleBee	46.6	83	67.15	30.82	<b>79.49</b>	<b>73.01</b>
LLaMA-7B	All	44.6	81	68.95	29.85	80.03	71.51
	Local	28.4	56	50.90	23.02	58.27	51.38
	Random + Local	28.0	63	51.26	21.76	53.94	49.30
	Attn Sinks + Local	41.6	<b>82</b>	58.12	27.40	78.07	67.80
	H2 + Local	41.4	78	63.54	27.50	77.31	65.82
	BumbleBee	<b>43.2</b>	79	<b>68.95</b>	27.74	78.24	<b>68.75</b>
	BumbleBee	<b>43.2</b>	79	63.90	<b>28.51</b>	<b>78.56</b>	68.19



# Experiments

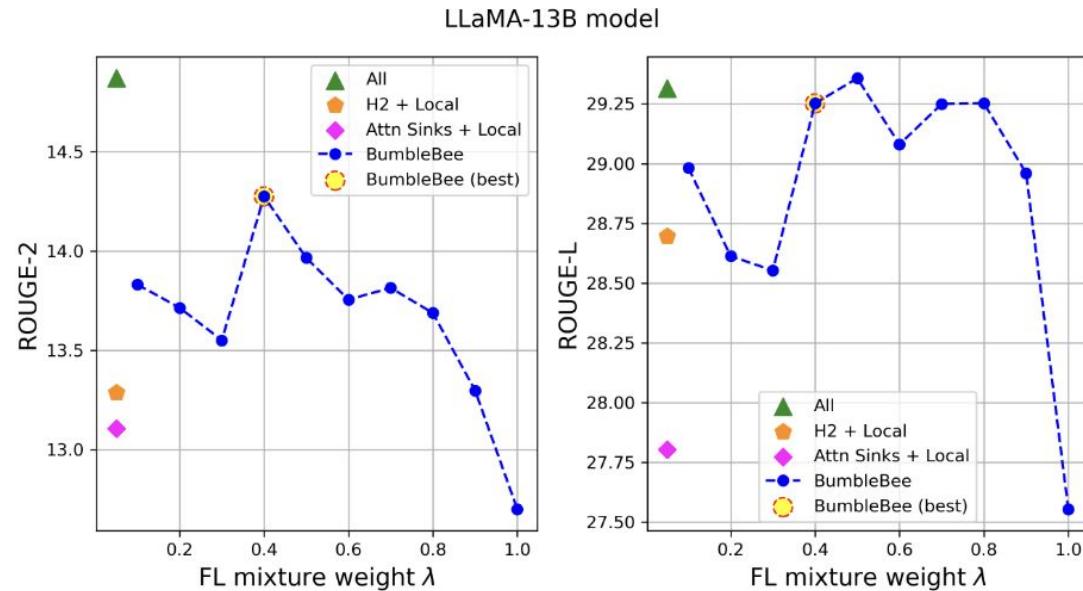
LongBench benchmark,  $\lambda=0.3$

Model	Method	Qasper	MultiFieldQA-en	HotpotQA	2WikiMQA	QMSum	TREC
LLaMA-7B-chat 4k	All*	19.20	36.80	25.40	32.80	20.80	61.5
	All (self)	21.60	36.76	27.55	31.58	20.78	64.0
	Attn Sinks + Local	14.74	22.93	22.08	29.73	19.25	56.0
	H2 (20%)	<b>19.82</b>	26.60	26.28	25.69	<b>21.45</b>	60.0
	BumbleBee (20%) ♡	19.37	27.73	26.14	27.67	20.68	<b>61.5</b>
	BumbleBee (20%) ♦	19.59	<b>28.60</b>	<b>28.99</b>	<b>30.19</b>	21.05	59.0
LongChat-7B 32k	H2 (SW, 20%)	21.64	30.72	14.07	15.10	18.11	40.5
	BumbleBee (SW, 20%) ♦	<b>23.27</b>	<b>33.16</b>	<b>22.52</b>	<b>17.58</b>	<b>20.27</b>	<b>44.5</b>

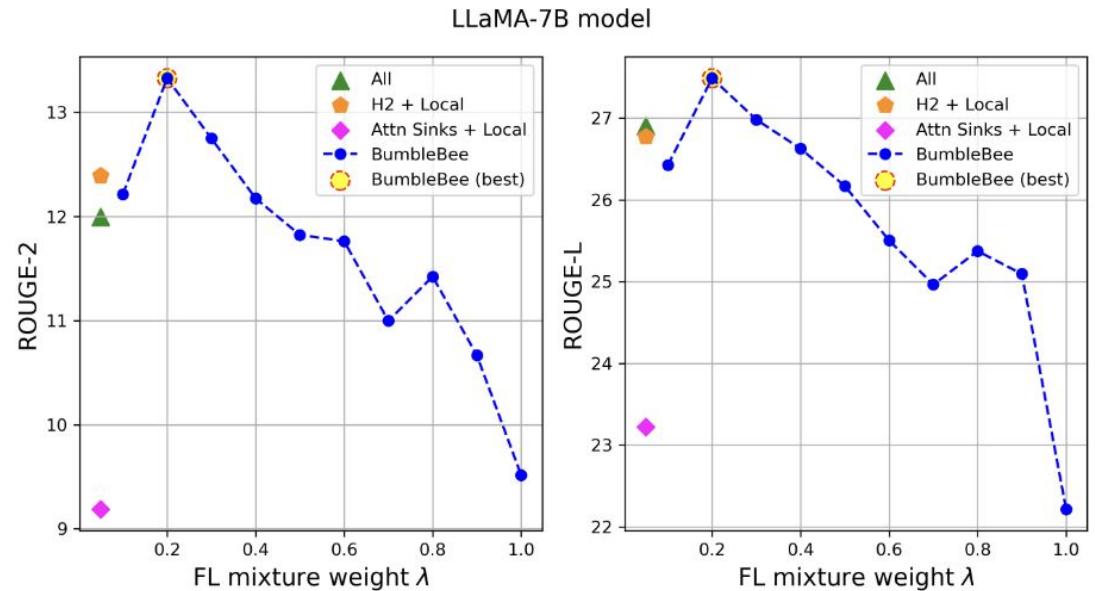
# Experiments

XSUM dataset, few-shot summarization task

$\lambda=0.2$  still better than H2+local (which equals to  $\lambda=0$  and  $\varphi$  is identity function)



(a) LLaMA-13B



(b) LLaMA-7B



# Experiments

---

Context reduction ratio	Original Context Length	
	16k	100k
1:1	59.30 ± 0.39	OOM
5:1	47.49 ± 4.16	71.50 ± 0.10
10:1	39.74 ± 1.31	48.16 ± 0.09

Table 6: Decoding speed (in ms/token) for two KV cache reduction ratios (5:1 and 10:1) and the baseline KV cache method using the entire context (1:1) across all heads. All experiments are performed on an A100 80GB GPU using the LongChat-7B-32k with a batch size of 1.

# Takeaways

---

- introduce diversity into selection aside from only importance(H2O)
- reframe eviction as subset selection problem, and submodularity guarantees that a simple Greedy Algorithm achieves a near-optimal solution

## Discussion:

- rely on submodular optimization tool(not open-source)
- redundancy/diversity matters
  - R-KV: heuristic ranking  $Z = \lambda \cdot \text{Imp} - (1-\lambda) \cdot \text{Redu}$
  - OmniKV: inter-layer redundancy

# Thank You!

jhu.edu

